

CSE 412 Database Management

Dr. Bharatesh Chakravarthi (BC)

Assistant Teaching Professor

School of Computing and Augmented Intelligence

Arizona State University

The Relational Model

What is a Relation?

- **Relation:** made up of 2 parts:
 - **Schema:** specifies name of relation, plus name and type of each column.
 - **Instance:** a table with rows and columns.
 - #rows = cardinality
 - #columns = arity

The diagram shows a relational table with three columns: Ssn, Name, and Address. The table has two rows of data: one row for 'smith' and one row for 'jones'. A red bracket on the left side indicates the cardinality of the table, which is 2. A red line above the table spans all three columns and is labeled 'arity' in red, indicating the number of columns in the relation.

Ssn	Name	Address
123	smith	main str
234	jones	forbes ave

Take-aways..

- *A database is a collection of **relations** (or **tables**)*
- *Each relation has a set of **attributes** (or **columns**)*
- *Each attribute has a name and a **domain** (or **type**)*
- *Each relation contains a set of **tuples** (or **rows**)*

What is a key?

- A set of attributes K is a key for a relation R if
 - In every instance of R , no two distinct tuples have the same values for all attributes in K
 - That is, K can serve as a “tuple identifier”
 - If K can serve as a “tuple identifier”, it is called as a Superkey
 - No proper subset of K satisfies the above condition
 - That is, K is minimal
 - Minimal superkey is called as a key
- Example:

Student (sid, name, login, age, gpa)

{sid, name}: superkey

{sid}: superkey, AND key

{name}: not superkey

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@cs	18	3.2
53650	Smith	smith@math	19	3.8

More examples..

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

- *Is name a key of User?*
 - Yes? Seems reasonable for this instance
 - No! User names are not unique *in general*

More examples..

<i>uid</i>	<i>name</i>	<i>age</i>	<i>pop</i>
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

- *Is name a key of User?*
 - Yes? Seems reasonable for this instance
 - No! User names are not unique *in general*
- *{uid} is the key*

More examples..

gid	name
abc	Book Club
gov	Student Government
dps	Dead Putting Society
...	...

- *Group (gid, name)*
- *Which set of attributes is the key?*

More examples..

<i>gid</i>	<i>name</i>
abc	Book Club
gov	Student Government
dps	Dead Putting Society
...	...

- *Group (gid, name)*
- *{gid} is the key*

More examples..

<i>uid</i>	<i>gid</i>
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

- Member (*uid*, *gid*)
- Which set of attributes is the key?

More examples..

<i>uid</i>	<i>gid</i>
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

- Member (*uid*, *gid*)
- {*uid*, *gid*} is the key
- A key can contain multiple attributes

Primary Keys..

- A *relation can have multiple keys!*
 - *Student: {ssn}, {student-id}, {driving-license#, state}*
 - *Employee: {ssn}, {employee-id}, {phone#}*
 - *computer: {mac-address}, {serial#}*
- *DBA (Database Admin) typically picks one as the “primary” key, and underline all its attributes*
- *Other keys are called candidate keys*

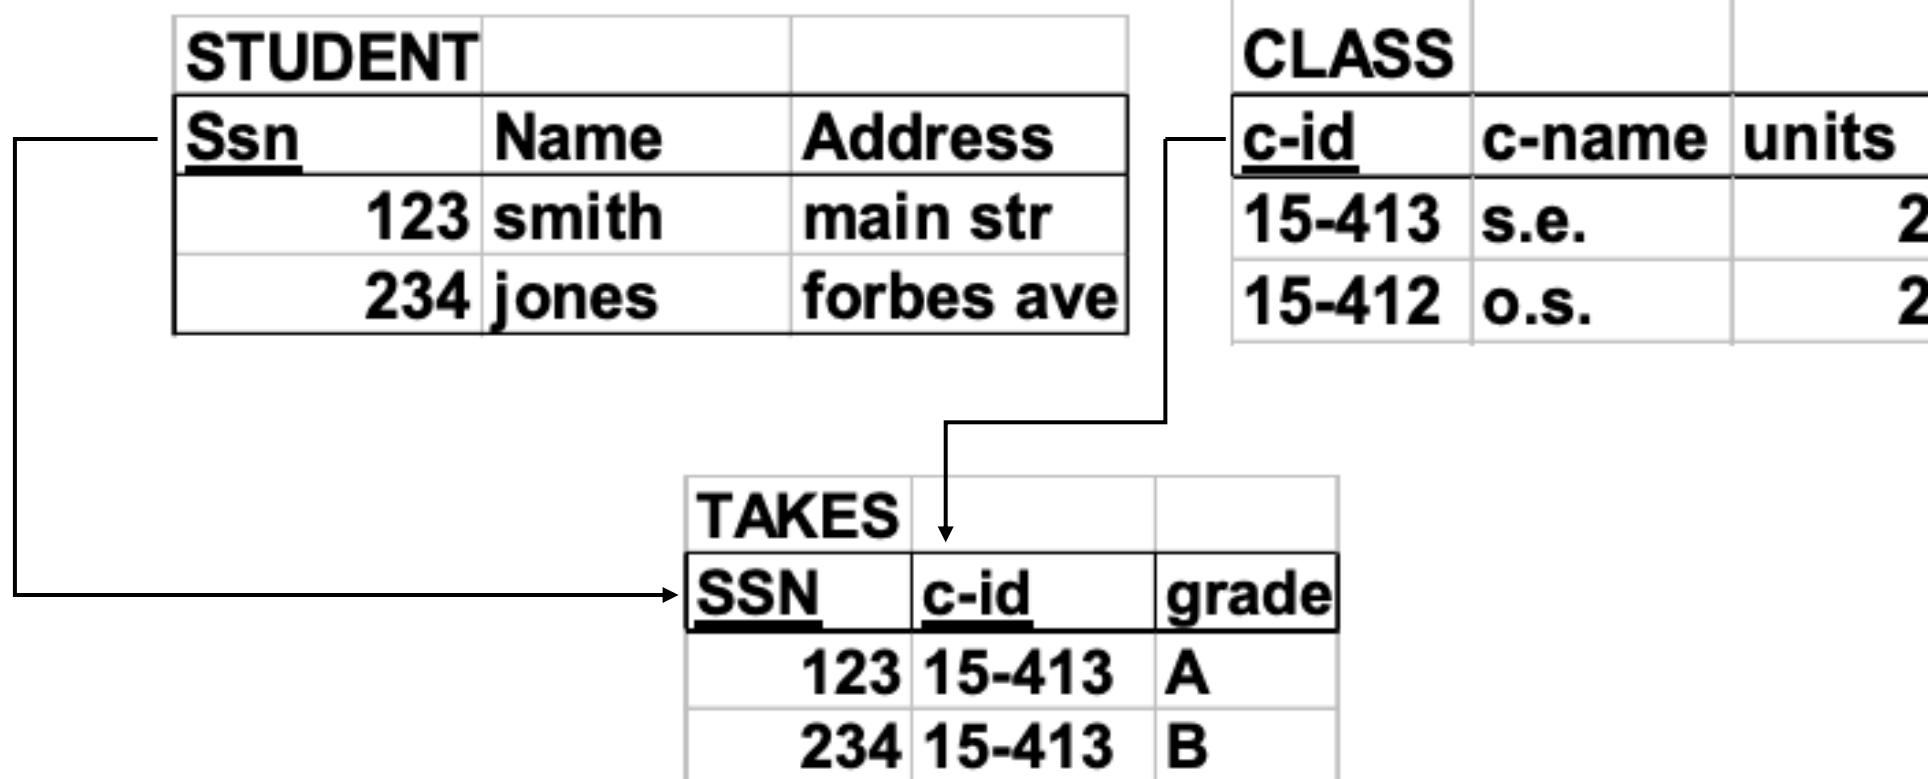
Primary Keys..

- *Example:*
 - *Address (street_address, city, state, zip)*
 - $\{street_address, city, state\}$
 - $\{street_address, zip\}$
 - *DBA typically picks one as the “primary” key, and underline all its attributes, e.g., Address (street address, city, state, zip)*
 - *Other keys are called candidate keys.*
e.g., {street_address, city, state}

What is a Database?

A database is a collection of *relations* (or *tables*)

Example 1



What is a Database?

A database is a collection of *relations* (or *tables*)

Example 2

User			
uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3

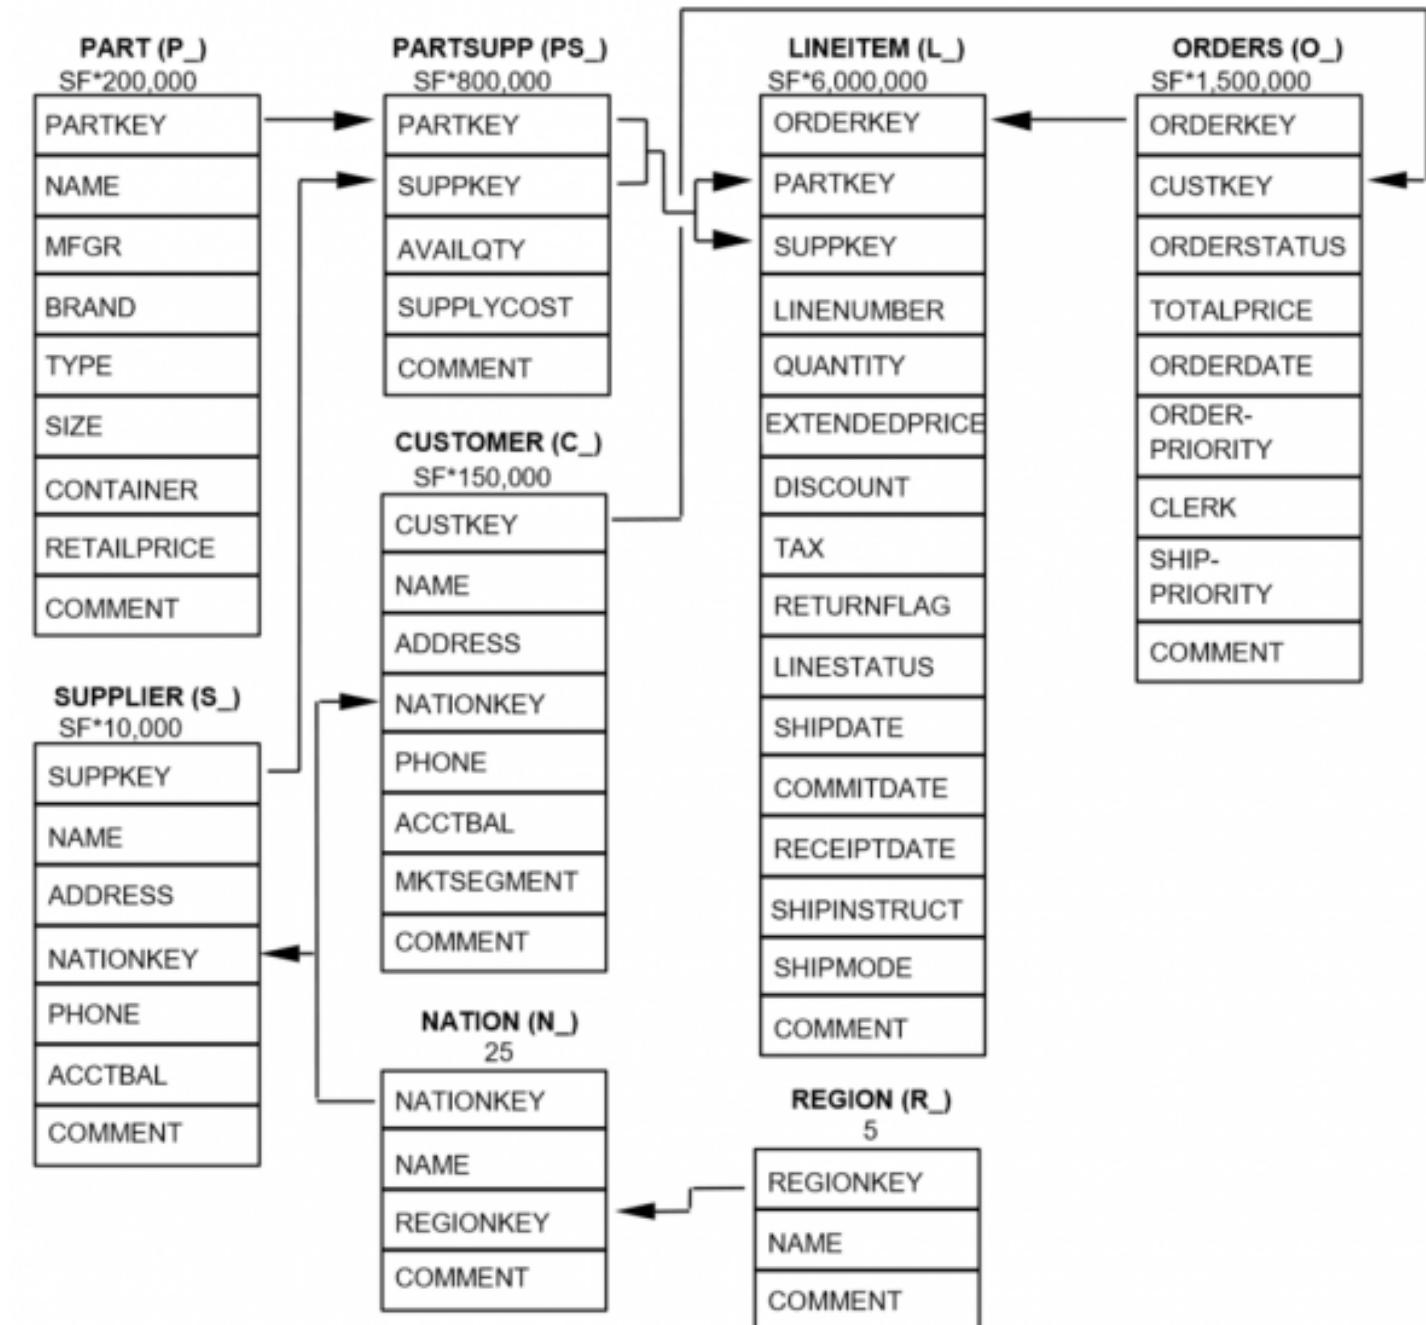
Group	
gid	name
abc	Book Club
gov	Student Government
dps	Dead Putting Society
...	...

Member	
uid	gid
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

What is a Database?

A database is a collection of relations (or tables)

Example 3



Foreign Keys, Referential Integrity

- *Foreign Key:* Set of attributes ‘referring’ to a tuple in another relation.
 - Must correspond to the primary key of the other relation
 - Like a ‘logical pointer’
- Foreign key constraints enforce *referential integrity* (i.e., no dangling references)

Data Definition Language (DDL)

Data Definition Language (DDL)

```
CREATE TABLE <table-name>(  
    [column-definition]*  
    [constraints]*  
    ) [table-options];
```

- **Column-Definition:** Comma separated list of column names with types.
- **Constraints:** Primary key, foreign key, and other meta-data attributes of columns.
- **Table-Options:** DBMS-specific options for the table (not SQL-92).

Data Definition Language (DDL)

```
CREATE TABLE student (
    sid    INT,
    name   VARCHAR(16),
    login  VARCHAR(32),
    age    SMALLINT,
    gpa    FLOAT
);

CREATE TABLE enrolled (
    sid    INT,
    cid    VARCHAR(32),
    grade  CHAR(1)
);
```

The diagram illustrates the data types defined in the DDL:

- Integer Range:** Points to the data types `INT`, `SMALLINT`, and `CHAR(1)`.
- Variable String Length:** Points to the data types `VARCHAR(16)`, `VARCHAR(32)`, and `VARCHAR(32)`.
- Fixed String Length:** Points to the data type `CHAR(1)`.

Common Data Types

- *CHAR(n) (fixed length), VARCHAR(n) (variable length)*
- *TINYINT (1 byte), SMALLINT (2 bytes), INT (4 bytes), , BIGINT (8 bytes),*
- *NUMERIC(p,d) (precision (no. of digits) and scale (digits to right of the decimal point)), FLOAT, DOUBLE, REAL*
- *DATE, TIME*
- *BINARY(n), VARBINARY(n), BLOB (image, video, files)*

Useful Non-standard Types

- *TEXT*
- *BOOLEAN*
- *ARRAY*
- *Geometric primitives*
- *XML/JSON*
- *Some systems also support user-defined types.*

Integrity Constraints

```
CREATE TABLE student (
    sid INT PRIMARY KEY,
    name VARCHAR(16),
    login VARCHAR(32) UNIQUE,
    age SMALLINT CHECK (age > 0),
    gpa FLOAT
);

CREATE TABLE enrolled (
    sid INT REFERENCES student (sid),
    cid VARCHAR(32) NOT NULL,
    grade CHAR(1),
    PRIMARY KEY (sid, cid)
);
```

PKey Definition

Type Attributes

FKey Definition

Primary Keys

- *Single-column primary key:*

```
CREATE TABLE student (
    sid    INT PRIMARY KEY,
    :)
```

- *Multi-column primary key:*

```
CREATE TABLE enrolled (
    :
    PRIMARY KEY (sid, cid)
```

Key declarations are part of the schema

CREATE TABLE enrolled

(sid CHAR(20),

cid CHAR(20),

grade CHAR(2),

PRIMARY KEY (*sid, cid*)

PRIMARY KEY == UNIQUE, NOT NULL

Foreign Key References

- *Single-column reference:*

```
CREATE TABLE enrolled (
    sid    INT REFERENCES student (sid),
    :)
```

- *Multi-column reference:*

```
CREATE TABLE enrolled (
    :
    FOREIGN KEY (sid, ...)
    REFERENCES student (sid, ...))
```

Foreign Keys in SQL

CREATE TABLE Enrolled (

sid CHAR(20),

cid CHAR(20),

grade CHAR(2),

PRIMARY KEY (sid,cid),

FOREIGN KEY (sid)

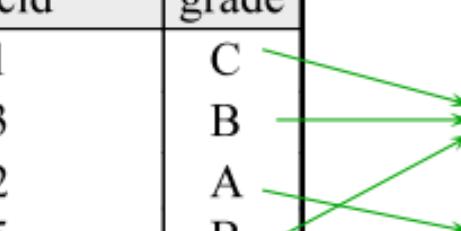
REFERENCES Students)

Enrolled

sid	cid	grade
53666	15-101	C
53666	18-203	B
53650	15-112	A
53666	15-105	B

Students

sid	name	login	age	gpa
53666	Jones	jones@cs	18	3.4
53688	Smith	smith@cs	18	3.2
53650	Smith	smith@math	19	3.8



Data Manipulation Language

Basic SQL

Basic Queries: SFW statements

- $\textcolor{red}{\text{SELECT } A_1, A_2, \dots, A_n}$
 $\textcolor{red}{\text{FROM } R_1, R_2, \dots, R_m}$
 $\textcolor{red}{\text{WHERE condition;}}$
- *Also called as SPJ (selection-projection-join) query - (selecting rows, projecting columns, and joining tables)* Select (σ), Project (π), Join (\bowtie)
- *Can correspond to (but not exactly equivalent to) relational algebra query:*
$$\pi_{A_1, A_2, \dots, A_n}(\sigma_{\text{condition}}(R_1 \times R_2 \times \dots \times R_m))$$

Example Database

User

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

Group

gid	name
abc	Book Club
gov	Student Government
dps	Dead Putting Society
...	...

Member

uid	gid
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

Example: reading a table

- ***SELECT * FROM User;***
 - *Single-table query, so no cross product here*
 - ***WHERE clause is optional***
 - **** is a short hand for “all columns”***

User

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

Example: selection and projection

- Name of users under 18

- *SELECT name
FROM User
WHERE age < 18*

- When was Lisa born?

- *SELECT 2025 - age
FROM User
WHERE name = 'Lisa';*

- *SELECT list can contain expressions*
 - *Can also use built-in functions such as SUBSTR, ABS, etc.*
 - *String literals (case sensitive) are enclosed in single quotes*

User

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

Example: join

- ID's and names of groups with a user whose name contains "Simpson"

User

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3

Group

gid	name
abc	Book Club
gov	Student Government
dps	Dead Putting Society
...	...

Member

uid	gid
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

Example: join

User

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

Group

gid	name
abc	Book Club
gov	Student Government
dps	Dead Putting Society
...	...

- *ID's and names of groups with a user whose name contains “Simpson”*
 - *SELECT Group.gid, Group.name
FROM User, Member, Group
WHERE User.uid = Member.uid
AND Member.gid = Group.gid
AND User.name LIKE '%Simpson%';*
 - *LIKE matches a string against a pattern*
 - *% matches any sequence of zero or more characters*
 - *Okay to omit table_name in table_name.column_name if column_name is unique*

Member

uid	gid
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

*Check out
TPC-H on PgAdmin*

