

CSE 412 Database Management

Dr. Bharatesh Chakravarthi (BC)


Assistant Teaching Professor

School of Computing and Augmented Intelligence

Arizona State University

Formal query languages

- How do we collect information?
 - Eg., find ssn's of students in CSE 412
- (Let's say: everything is a set!)
- One solution: Rel. algebra, ie., set operators
 - Q1: Which ones??
 - Q2: What is a minimal set of operators?



Edgar F. Codd proposed relational model and Relational algebra in 1970s

Relational Operators

Fundamental Operators

- Selection: $\sigma_p R$
- Projection: $\pi_L R$
- Cross product: $R \times S$
- Union: $R \cup S$
- Difference: $R - S$

Derived Operators

- Intersection: $R \cap S$
- Join: $R \bowtie_p S$
- Natural join: $R \bowtie S$
- Renaming: $\rho_{S(A_1, A_2, \dots)} R$

Relational Operators

Fundamental Operators

- Selection: $\sigma_p R$
- Projection: $\pi_L R$
- Cross product: $R \times S$
- Union: $R \cup S$
- Difference: $R - S$

Derived Operators

- Intersection: $R \cap S$
- Join: $R \bowtie_p S$
- Natural join: $R \bowtie S$
- Renaming: $\rho_{S(A_1, A_2, \dots)} R$


Compose operators to make complex queries.

Selection Example

- Users with popularity higher than 0.5

$$\sigma_{pop>0.5} User$$

<i>uid</i>	<i>name</i>	<i>age</i>	<i>pop</i>
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...




<i>uid</i>	<i>name</i>	<i>age</i>	<i>pop</i>
142	Bart	10	0.9
857	Lisa	8	0.7
...

Selection Example

- Users with popularity higher than 0.5

$$\sigma_{pop>0.5} User$$

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...



uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

Corresponding SQL:
Select * FROM User WHERE pop > 0.5

Selection

- Input: a table R
- Notation: $\sigma_p R$
 - P is called as a *selection condition* (or *predicate*)
- **Purpose:** filter rows according to some criteria
- Output: same columns as R, but only rows of R that satisfy p

More on Selection

- Selection condition can include any column of R,
 - constants, comparison ($=$, \leq , etc.) and
 - Boolean connectives (\wedge : and, \vee : or, \neg : not)
- Example: users with popularity at least 0.9 **and** age under 10
or above 12

$$\sigma_{pop \geq 0.9 \wedge (age < 10 \vee age > 12)} User$$

More on Selection

- Selection condition can include any column of R,
 - constants, comparison ($=$, \leq , etc.) and
 - Boolean connectives (\wedge : and, \vee : or, \neg : not)
- Example: users with popularity at least 0.9 and age under 10 or above 12

$$\sigma_{pop \geq 0.9 \wedge (age < 10 \vee age > 12)} User$$

Corresponding SQL:

Select * FROM User WHERE pop > 0.5 AND
(age < 10 OR age > 12)

More on Selection

- You must be able to evaluate the condition over each single row of the input table!
- Example: the most popular user

$\sigma_{pop \geq \text{every } pop \text{ in User}}$ User
WRONG!

Correct Form in SQL could be:

```
SELECT *  
FROM User  
WHERE pop = (SELECT MAX(pop) FROM User)
```

Relational Operators

Fundamental Operators

- Selection: $\sigma_p R$
- Projection: $\pi_L R$
- Cross product: $R \times S$
- Union: $R \cup S$
- Difference: $R - S$

Derived Operators

- Intersection: $R \cap S$
- Join: $R \bowtie_p S$
- Natural join: $R \bowtie S$
- Renaming: $\rho_{S(A_1, A_2, \dots)} R$


Compose operators to make complex queries.

Projection Examples

- IDs and names of all users

$$\pi_{uid, name} \text{User}$$

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...




uid	name
142	Bart
123	Milhouse
857	Lisa
456	Ralph
...	...

Projection Examples

- IDs and names of all users

$\pi_{uid, name} User$

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...




uid	name
142	Bart
123	Milhouse
857	Lisa
456	Ralph
...	...

Corresponding SQL:
Select UID NAME
FROM User

More on projection

- Duplicate output rows are removed (by definition)
- Example: user ages

$\pi_{age} User$




More on projection

Corresponding SQL:
**Select DISTINCT
AGE FROM User**

- Duplicate output rows are removed (by definition)
- Example: user ages

$\pi_{age} User$

<i>uid</i>	<i>name</i>	<i>age</i>	<i>pop</i>
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...



Projection

- Input: a table R
- Notation: $\pi_L R$
 - L is a list of columns in R
- **Purpose:** output chosen columns
- Output: same rows, but only the columns in L

Test

- What does this mean:

$\sigma_{pop > 0.5} (\pi_{name, age}) \text{User}$

User

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

Test

- How to translate this to SQL?

$\sigma_{pop > 0.5} (\pi_{name, age}) User$

User

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

Test

- How to translate this to SQL?

$\sigma_{pop > 0.5} (\pi_{name, age}) User$

User

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

```
SELECT Name, Age  
FROM User  
WHERE pop > 0.5
```

Relational Operators

Fundamental Operators

- ~~Selection:~~ $\sigma_p R$
- ~~Projection:~~ $\pi_L R$
- Cross product: $R \times S$
- Union: $R \cup S$
- Difference: $R - S$

Derived Operators

- Intersection: $R \cap S$
- Join: $R \bowtie_p S$
- Natural join: $R \bowtie S$
- Renaming: $\rho_{S(A_1, A_2, \dots)} R$

Compose operators to make complex queries.

Cross product example

User×Member

<i>uid</i>	<i>name</i>	<i>age</i>	<i>pop</i>
123	Milhouse	10	0.2
857	Lisa	8	0.7
...

<i>uid</i>	<i>gid</i>
123	gov
857	abc
857	gov
...	...

<i>uid</i>	<i>name</i>	<i>age</i>	<i>pop</i>	<i>uid</i>	<i>gid</i>
123	Milhouse	10	0.2	123	gov
123	Milhouse	10	0.2	857	abc
123	Milhouse	10	0.2	857	gov
857	Lisa	8	0.7	123	gov
857	Lisa	8	0.7	857	abc
857	Lisa	8	0.7	857	gov
...

Cross product example

User×Member

<i>uid</i>	<i>name</i>	<i>age</i>	<i>pop</i>
123	Milhouse	10	0.2
857	Lisa	8	0.7
...

<i>uid</i>	<i>gid</i>
123	gov
857	abc
857	gov
...	...

<i>uid</i>	<i>name</i>	<i>age</i>	<i>pop</i>	<i>uid</i>	<i>gid</i>
123	Milhouse	10	0.2	123	gov
123	Milhouse	10	0.2	857	abc
123	Milhouse	10	0.2	857	gov
857	Lisa	8	0.7	123	gov
857	Lisa	8	0.7	857	abc
857	Lisa	8	0.7	857	gov
...

Corresponding SQL:
Select *
FROM User, Member

Cross product

- Input: two tables R and S
- Notation: $R \times S$
- **Purpose: pairs of rows from two tables**
- Output: for each row r in R and each s in S, output a row rs (concatenation of r and s)

A note on column/row ordering

- Ordering of columns/rows is unimportant as far as contents are concerned

<i>uid</i>	<i>name</i>	<i>age</i>	<i>pop</i>	<i>uid</i>	<i>gid</i>	<i>uid</i>	<i>gid</i>	<i>uid</i>	<i>name</i>	<i>age</i>	<i>pop</i>
123	Milhouse	10	0.2	123	gov	123	gov	123	Milhouse	10	0.2
123	Milhouse	10	0.2	857	abc	857	abc	123	Milhouse	10	0.2
123	Milhouse	10	0.2	857	gov	857	gov	123	Milhouse	10	0.2
857	Lisa	8	0.7	123	gov	123	gov	857	Lisa	8	0.7
857	Lisa	8	0.7	857	abc	857	abc	857	Lisa	8	0.7
857	Lisa	8	0.7	857	gov	857	gov	857	Lisa	8	0.7
...

- So cross product is **commutative**, i.e., for any R and S, RxS = SxR (up to the ordering of columns and rows)

Relational Operators

Fundamental Operators

- Selection: $\sigma_p R$
- Projection: $\pi_L R$
- Cross product: $R \times S$
- Union: $R \cup S$
- Difference: $R - S$

Derived Operators

- Intersection: $R \cap S$
- Join: $R \bowtie_p S$
- Natural join: $R \bowtie S$
- Renaming: $\rho_{S(A_1, A_2, \dots)} R$

Compose operators to make complex queries.

Test

- What does this mean:


$\sigma_{User.uid = Member.uid}$ (*User x Member*)

Test

- What does this mean:

$\sigma_{User.uid = Member.uid}$ (*User x Member*)

User x Member



uid	name	age	pop	uid	gid
123	Milhouse	10	0.2	123	gov
123	Milhouse	10	0.2	857	abc
123	Milhouse	10	0.2	857	gov
857	Lisa	8	0.7	123	gov
857	Lisa	8	0.7	857	abc
857	Lisa	8	0.7	857	gov
...

Test

- What does this mean:

$\sigma_{User.uid = Member.uid}$ (*User x Member*)

User x Member

uid	name	age	pop
123	Milhouse	10	0.2
857	Lisa	8	0.7
...

uid	gid
123	gov
857	abc
857	gov
...	...

uid	name	age	pop	uid	gid
123	Milhouse	10	0.2	123	gov
857	Lisa	8	0.7	857	abc
857	Lisa	8	0.7	857	gov
...

$\sigma_{User.uid = Member.uid}$

uid	name	age	pop	uid	gid
123	Milhouse	10	0.2	123	gov
123	Milhouse	10	0.2	857	abc
123	Milhouse	10	0.2	857	gov
857	Lisa	8	0.7	123	gov
857	Lisa	8	0.7	857	abc
857	Lisa	8	0.7	857	gov
...

Test

- What does this mean:

$\sigma_{User.uid = Member.uid}$ (*User x Member*)

It returns you the
group information
for each user

uid	name	age	pop	uid	gid
123	Milhouse	10	0.2	123	gov
857	Lisa	8	0.7	857	abc
857	Lisa	8	0.7	857	gov
...

uid	name	age	pop
123	Milhouse	10	0.2
857	Lisa	8	0.7
...


User x Member

uid	gid
123	gov
857	abc
857	gov
...	...

$\sigma_{User.uid = Member.uid}$

uid	name	age	pop	uid	gid
123	Milhouse	10	0.2	123	gov
123	Milhouse	10	0.2	857	abc
123	Milhouse	10	0.2	857	gov
857	Lisa	8	0.7	123	gov
857	Lisa	8	0.7	857	abc
857	Lisa	8	0.7	857	gov
...

Because this query pattern is so popular, a new operator is derived from cross product (\times) and selection (σ), *called join* (\bowtie)


$$\sigma_{User.uid = Member.uid} (User \times Member)$$

$$User \bowtie_{User.uid = Member.uid} Member$$

Derived operator: Join

- (a.k.a. “theta-join”)
- Input: two tables R and S
- Notation: $R \bowtie_p S$
 - p is called a join condition (or predicate)
- **Purpose:** relate rows from two tables according to some criteria
- Output: for each row r in R, and each rows s in S, output rs , if r and s satisfy p
- *Shorthand for $\sigma_p(R \times S)$*

Join example


- Info about users, plus IDs of their groups



Corresponding SQL 1:

**Select * FROM User, Member WHERE User.uid =
Member.uid**

$\sigma_{User.uid = Member.uid} (User \times Member)$



$User \bowtie_{User.uid = Member.uid} Member$

Corresponding SQL 2:


**Select * FROM User JOIN Member ON User.uid =
Member.uid**

Corresponding SQL 1:

Select * FROM User, Member WHERE User.uid =
Member.uid

$\sigma_{User.uid = Member.uid} (User \times Member)$

$User \bowtie_{User.uid = Member.uid} Member$



Both
Correct!

Corresponding SQL 2:

Select * FROM User JOIN Member ON User.uid =
Member.uid

Relational Operators

Fundamental Operators

- Selection: $\sigma_p R$
- Projection: $\pi_L R$
- Cross product: $R \times S$
- Union: $R \cup S$
- Difference: $R - S$

Derived Operators

- Intersection: $R \cap S$
- Join: $R \bowtie_p S$
- Natural join: $R \bowtie S$
- Renaming: $\rho_{S(A_1, A_2, \dots)} R$

Compose operators to make complex queries.