

Relational Operators

Fundamental Operators

- ~~Selection:~~ $\sigma_p R$
- ~~Projection:~~ $\pi_L R$
- ~~Cross product:~~ $R \times S$
- Union: $R \cup S$
- Difference: $R - S$

Derived Operators

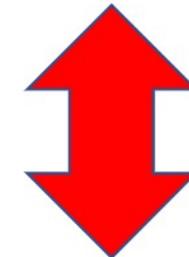
- Intersection: $R \cap S$
- ~~Join:~~ $R \bowtie_p S$
- Natural join: $R \bowtie S$
- Renaming: $\rho_{S(A_1, A_2, \dots)} R$

Compose operators to make complex queries.

- New Problem:
 - The User table and the Member table share one column (UID)
 - Can we simplify the join predicate, if two tables join on all shared columns?

User				Member	
uid	name	age	po	uid	gid
142	Bart	10	0	142	dps
123	Milhouse	10	0	123	gov
857	Lisa	8	0	857	abc
456	Ralph	8	0	857	gov
...	456	abc
				456	gov
			

User $\bowtie_{User.uid = Member.uid}$ Member



User \bowtie Member

Derived operator: Natural Join

- Input: two tables R and S
- Notation: $R \bowtie S$
- Purpose: relate rows from two tables
 - Enforce equality between identically named columns
 - Eliminate one copy of identically named columns
- Shorthand for $\pi_L(R \bowtie_p S)$, where
 - p equates each pair of columns common to R and S
 - L is the union of column names from R and S (with duplicate columns)

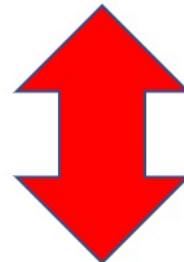
User

uid	name	age	po
142	Bart	10	0
123	Milhouse	10	0
857	Lisa	8	0
456	Ralph	8	0
...

Member

uid	gid
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

User \bowtie User.uid = Member.uid Member



User \bowtie Member

Corresponding SQL:

Select * FROM User, Member WHERE User.uid = Member.uid

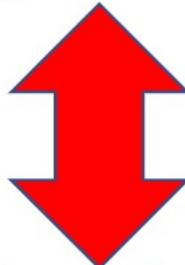
Select * FROM User JOIN Member ON User.uid = Member.uid

Select * FROM User NATURAL JOIN Member

Corresponding SQL 1:

Select * FROM User, Member WHERE User.uid =
Member.uid

$\sigma_{User.uid = Member.uid} (User \times Member)$



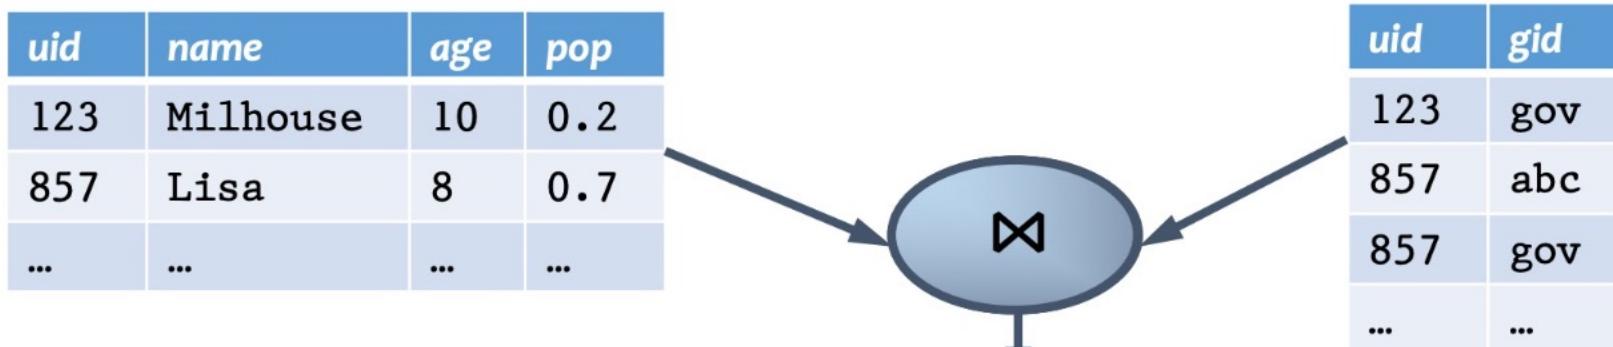
User $\bowtie_{User.uid = Member.uid}$ Member

Corresponding SQL 2:

Select * FROM User JOIN Member ON User.uid =
Member.uid

Natural Join Example

$$\begin{aligned}User \bowtie Member &= \pi_?(User \bowtie_? Member) \\&= \pi_{uid, name, age, pop, gid} (User \bowtie User.uid = Member.Member.uid)\end{aligned}$$



uid	name	age	pop		gid
123	Milhouse	10	0.2		gov
857	Lisa	8	0.7		abc
857	Lisa	8	0.7		gov
...

Relational Operators

Fundamental Operators

- Selection: $\sigma_p R$
- Projection: $\pi_L R$
- Cross product: $R \times S$
- Union: $R \cup S$
- Difference: $R - S$

Derived Operators

- Intersection: $R \cap S$
- Join: $R \bowtie_p S$
- Natural join: $R \bowtie S$
- Renaming: $\rho_{S(A_1, A_2, \dots)} R$

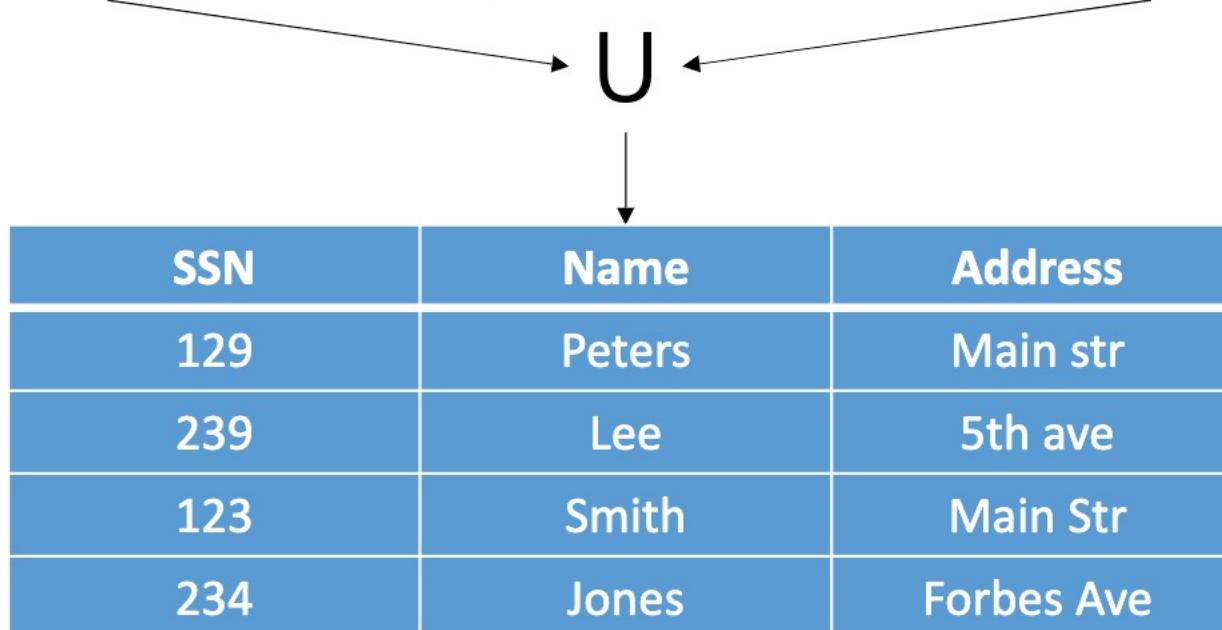
Compose operators to make complex queries.

Union Example

- Q: find all students (part or full time)
- A: PT-STUDENT union FT-STUDENT

FT-STUDENT		
Ssn	Name	Address
129	peters	main str
239	lee	5th ave

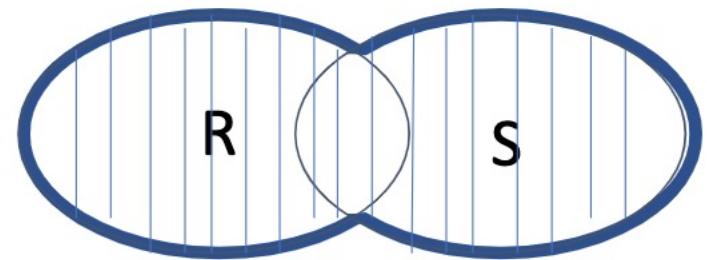
PT-STUDENT		
Ssn	Name	Address
123	smith	main str
234	jones	forbes ave



Two tables are ‘union compatible’ if they have the same attributes (‘domains’)

Union

- Input: two tables R and S
- Notation: $R \cup S$
 - R and S must have identical schema
- Output:
 - Has the same schema as R and S
 - Contains all rows in R and all rows in S (with duplicate rows removed)



Relational Operators

Fundamental Operators

- Selection: $\sigma_p R$
- Projection: $\pi_L R$
- Cross product: $R \times S$
- Union: $R \cup S$
- Difference: $R - S$

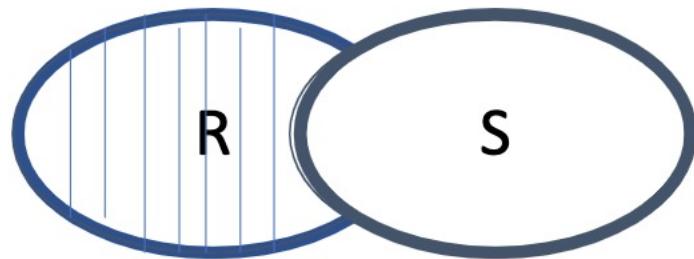
Derived Operators

- Intersection: $R \cap S$
- Join: $R \bowtie_p S$
- Natural join: $R \bowtie S$
- Renaming: $\rho_{S(A_1, A_2, \dots)} R$

Compose operators to make complex queries.

Difference

- Input: two tables R and S
- Notation: $R - S$
 - R and S must have identical schema
- Output:
 - Has the same schema as R and S
 - Contains all rows in R that are not in S



Difference Example

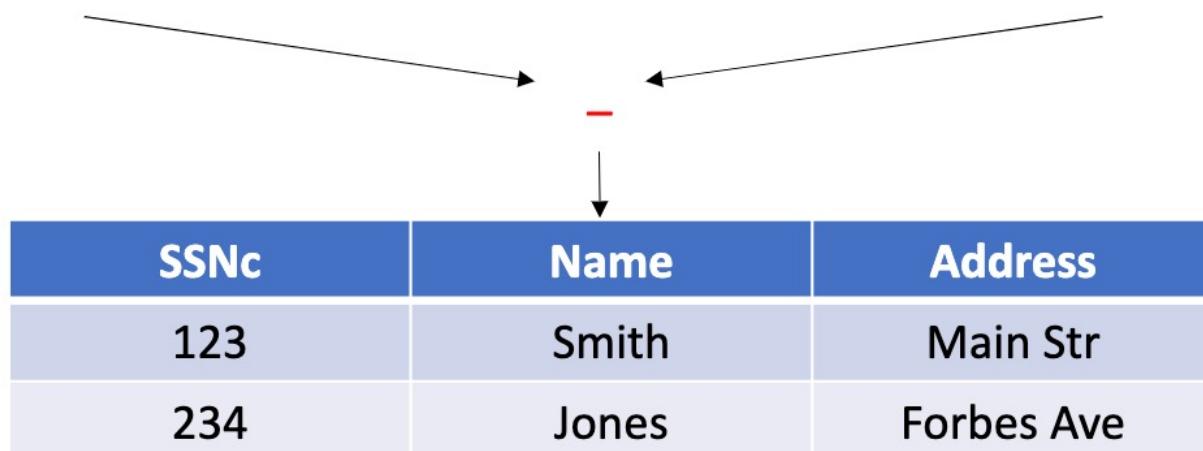
- Q: find all students who is not a staff
- A: Student - Staff

Student

SSNc	Name	Address
129	Peters	Main str
239	Lee	5th ave
123	Smith	Main Str
234	Jones	Forbes Ave

Staff

SSN	Name	Address
129	Peters	Main str
239	Lee	5th ave
323	Sarah	Sunset Blvd
432	Mary	Sunset Blvd



Relational Operators

Fundamental Operators

- Selection: $\sigma_p R$
- Projection: $\pi_L R$
- Cross product: $R \times S$
- Union: $R \cup S$
- Difference: $R - S$

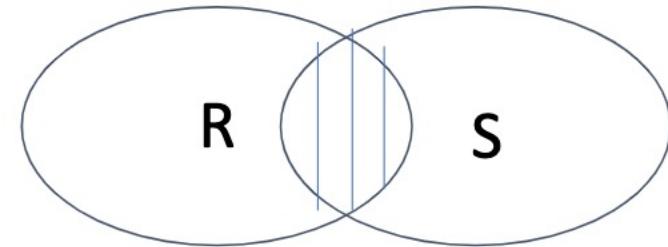
Derived Operators

- Intersection: $R \cap S$
- Join: $R \bowtie_p S$
- Natural join: $R \bowtie S$
- Renaming: $\rho_{S(A_1, A_2, \dots)} R$

Compose operators to make complex queries.

Derived Operator: Intersection

- Input: two tables R and S
- Notation: $R \cap S$
 - R and S must have identical schemas
- Output:
 - Has the same schema as R and S
 - Contains all rows that are in both relations
- Shorthand for $R - (R - S)$
- Also equivalent to $S - (S - R)$
- And to $R \bowtie S$



Intersection Example

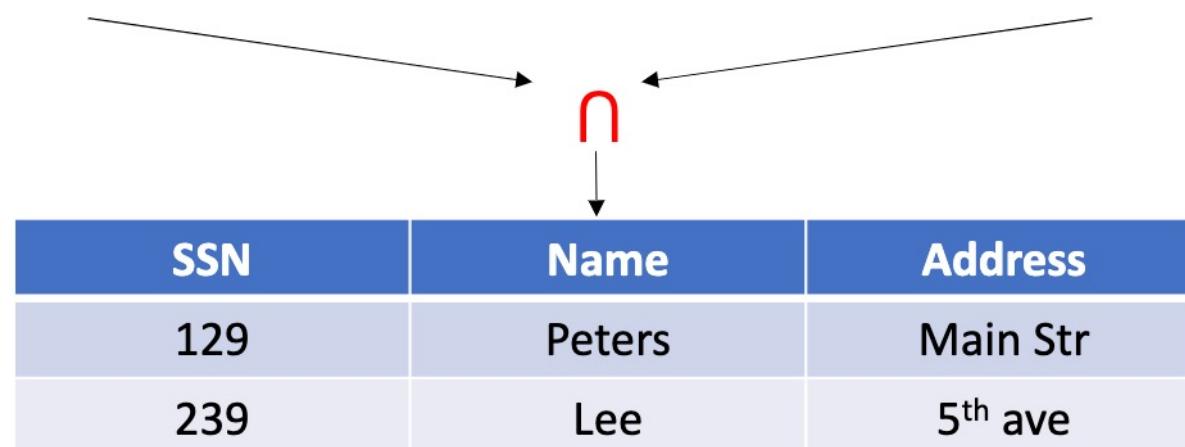
- Student intersection Staff
- Student – (Student – Staff)

Student

SSN	Name	Address
129	Peters	Main str
239	Lee	5th ave
123	Smith	Main Str
234	Jones	Forbes Ave

Staff

SSN	Name	Address
129	Peters	Main str
239	Lee	5th ave
323	Sarah	Sunset Blvd
432	Mary	Sunset Blvd



Relational Operators

Fundamental Operators

- ~~Selection:~~ $\sigma_p R$
- ~~Projection:~~ $\pi_L R$
- ~~Cross product:~~ $R \times S$
- ~~Union:~~ $R \cup S$
- ~~Difference:~~ $R - S$

Derived Operators

- ~~Intersection:~~ $R \cap S$
- ~~Join:~~ $R \bowtie_p S$
- ~~Natural join:~~ $R \bowtie S$
- Renaming: $\rho_{S(A_1, A_2, \dots)} R$

Compose operators to make complex queries.

Renaming

- Input: a table R
- Notation: $\rho_s R$, $\rho_{(A_1, A_2, \dots)} R$, or $\rho_{S(A_1, A_2, \dots)} R$
- Purpose: “rename” a table or its columns
- Output: a table with the same rows as R, but called differently
- Used to
 - Avoid confusion caused by identical column names
 - Create identical column names for natural joins
- It doesn’t modify the database
 - Think of the renamed table as a copy of the original

Renaming Example

- IDs of users who belong to at least two groups

Member $\bowtie_?$ *Member*

$\pi_{uid} \left(\text{Member} \bowtie_{\text{Member}.uid = \text{Member}.uid \wedge \text{Member}.gid \neq \text{Member}.gid} \text{Member} \right)$

WRONG!

$\pi_{uid_1} \left(\begin{array}{l} \rho_{(uid_1, gid_1)} \text{Member} \\ \bowtie_{uid_1 = uid_2 \wedge gid_1 \neq gid_2} \\ \rho_{(uid_2, gid_2)} \text{Member} \end{array} \right)$

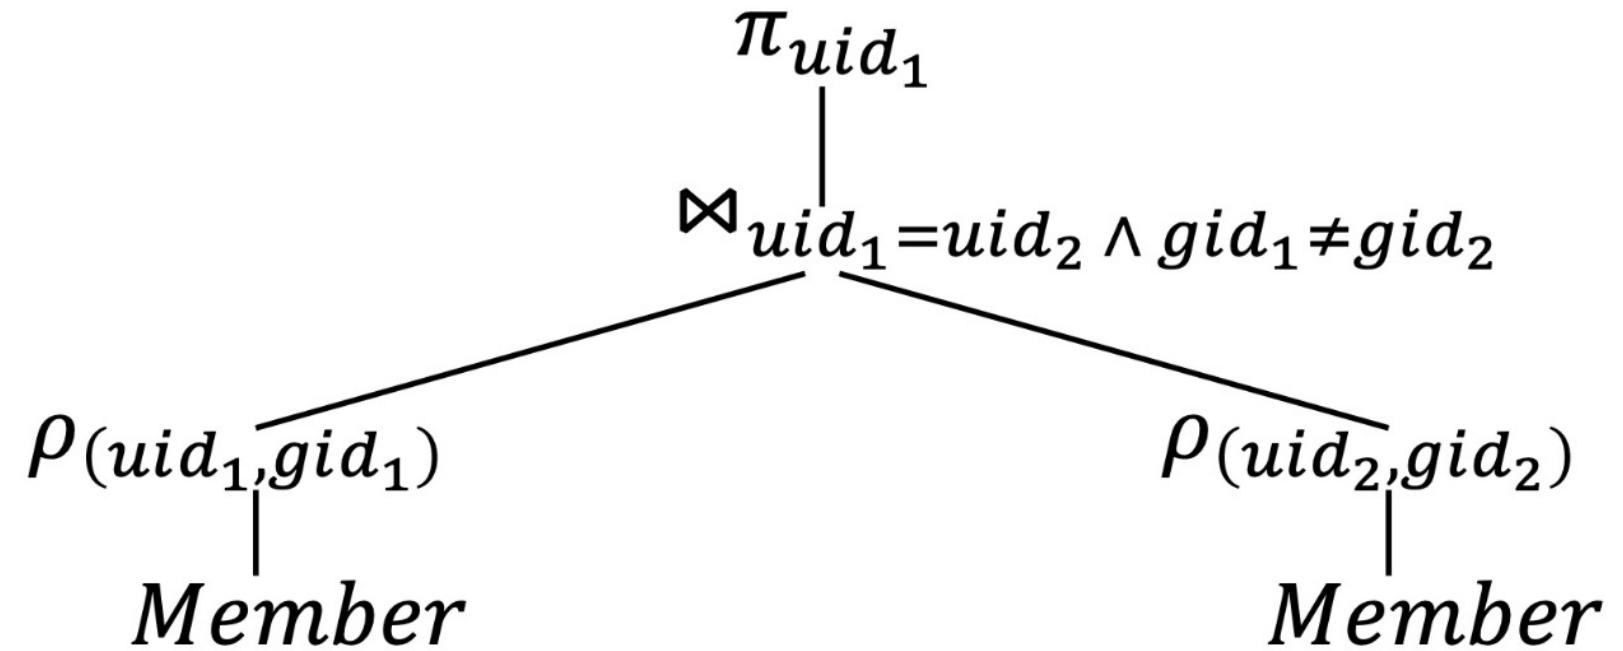
Member

<i>uid</i>	<i>gid</i>
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

Expression Tree Notation

Member

	<i>uid</i>	<i>gid</i>
	142	dps
	123	gov
	857	abc
	857	gov
	456	abc
	456	gov



Relational Operators

Fundamental Operators

- ~~Selection:~~ $\sigma_p R$
- ~~Projection:~~ $\pi_L R$
- ~~Cross product:~~ $R \times S$
- ~~Union:~~ $R \cup S$
- ~~Difference:~~ $R - S$

Derived Operators

- ~~Intersection:~~ $R \cap S$
- ~~Join:~~ $R \bowtie_p S$
- ~~Natural join:~~ $R \bowtie S$
- ~~Renaming:~~ $\rho_{S(A_1, A_2, \dots)} R$

Compose operators to make complex queries.

An exercise: Names of users in Lisa's groups

User

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

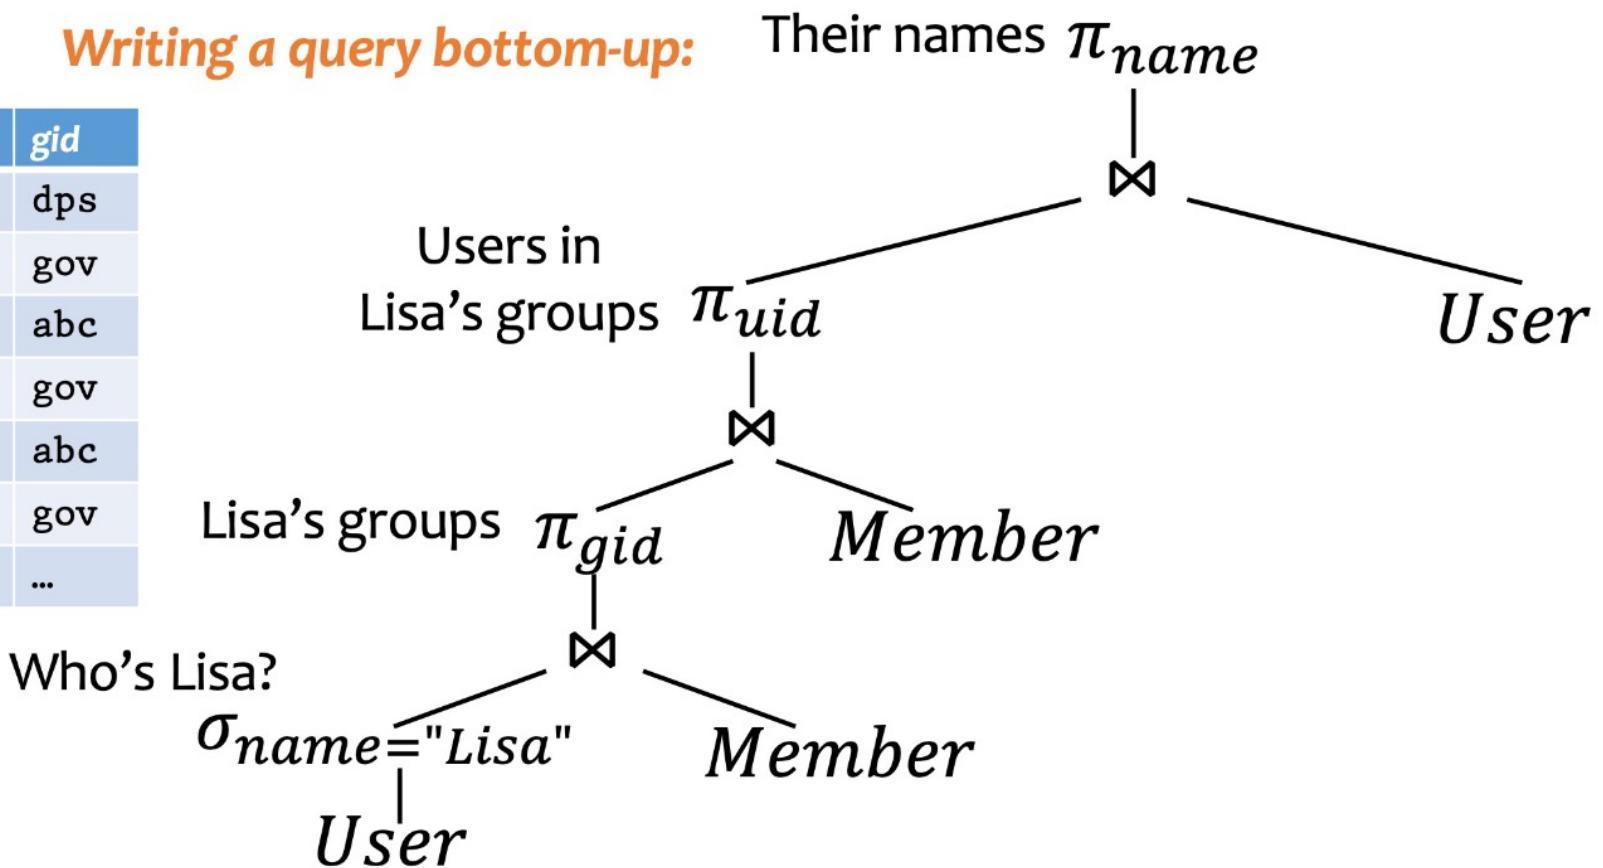
Group

gid	name
abc	Book Club
gov	Student Government
dps	Dead Putting Society
...	...

Member

uid	gid
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

Writing a query bottom-up:



Another exercise: IDs of groups that Lisa doesn't belong to

Writing a query top-down:

User

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

Group

gid	name
abc	Book Club
gov	Student Government
dps	Dead Putting Society
...	...

Member

uid	gid
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

All group IDs

π_{gid}

Group

IDs of Lisa's groups

π_{gid}

Member $\sigma_{name='Lisa'}$

User

More Examples

STUDENT		
Ssn	Name	Address
123	smith	main str
234	jones	forbes ave

CLASS		
c-id	c-name	units
413	s.e.	2
412	o.s.	2

TAKES		
SSN	c-id	grade
123	413	A
234	413	B

- Find course names of 'smith'
- Find ssn of 'overworked' students, ie., who take 412, 413, 415

Find course names of ‘Smith’

$$\pi_{c-name} [\sigma_{name='smith'} ($$

STUDENT \bowtie *TAKES* \bowtie *CLASS*

)]



A diagram illustrating the scope of the projection in the query. Two horizontal arrows point to the right from the closing bracket of the query. The top arrow is blue and points to the closing bracket of the relationship name 'TAKES'. The bottom arrow is orange and also points to the same closing bracket, indicating that the projection applies to all attributes of the 'TAKES' relationship.

Find ssn of ‘overworked’ students, ie., who take 412, 413, 415

$$\sigma_{c-name=412}(TAKES) \cap$$

$$\sigma_{c-name=413}(TAKES) \cap$$

$$\sigma_{c-name=415}(TAKES)$$

