

CSE 412 Database Management

Dr. Bharatesh Chakravarthi (BC)

Assistant Teaching Professor

School of Computing and Augmented Intelligence

Arizona State University

Review of SQL Basics

Review of Data Definition Language

```
CREATE TABLE <table-name>(  
    [column-definition]*  
    [constraints]*  
) [table-options];
```

- **Column-Definition:** Comma separated list of column names with types.
- **Constraints:** Primary key, foreign key, and other meta-data attributes of columns.
- **Table-Options:** DBMS-specific options for the table (not SQL-92).

Data Definition Language (DDL)

```
CREATE TABLE student (
    sid    INT,  

    name   VARCHAR(16),  

    login  VARCHAR(32),  

    age    SMALLINT,  

    gpa    FLOAT  

);  
  
CREATE TABLE enrolled (  

    sid    INT,  

    cid    VARCHAR(32),  

    grade  CHAR(1)  

);
```

The diagram illustrates the data types defined in the DDL. Red arrows point from specific column definitions to descriptive labels:

- An arrow points from the `sid` definition in the `student` table to the label **Integer Range**.
- An arrow points from the `name`, `login`, and `grade` definitions to the label **Variable String Length**.
- An arrow points from the `cid` definition in the `enrolled` table to the label **Fixed String Length**.

Integrity Constraints

```
CREATE TABLE student (
    sid    INT PRIMARY KEY,
    name   VARCHAR(16),
    login  VARCHAR(32) UNIQUE,
    age    SMALLINT CHECK (age > 0),
    gpa    FLOAT
);
CREATE TABLE enrolled (
    sid    INT REFERENCES student (sid),
    cid    VARCHAR(32) NOT NULL,
    grade  CHAR(1),
    PRIMARY KEY (sid, cid)
);
```

PKey Definition

Type Attributes

FKey Definition

Primary Keys

- Single-column primary key:

```
CREATE TABLE student (
    sid    INT PRIMARY KEY,
    :)
```

- Multi-column primary key:

```
CREATE TABLE enrolled (
    :
    PRIMARY KEY (sid, cid)
```

Foreign Key References

- Single-column reference:

```
CREATE TABLE enrolled (
    sid    INT REFERENCES student (sid),
    :)
```

- Multi-column reference:

```
CREATE TABLE enrolled (
    :
    FOREIGN KEY (sid, ...)
    REFERENCES student (sid, ...))
```

Value Constraints

- Ensure one-and-only-one value exists:

```
CREATE TABLE student (
    login VARCHAR(32) UNIQUE,
```

- Make sure a value is not null:

```
CREATE TABLE enrolled (
    cid    VARCHAR(32) NOT NULL,
```

Tuple- and Attribute- based Checks

- Associated with a *single table*
- Only checked *when a tuple/attribute is inserted/updated*
 - Reject if the condition evaluates to FALSE
 - TRUE and UNKNOWN are fine
- Examples:
 - CREATE TABLE User(... age INTEGER CHECK(age IS NULL OR age > 0), ...);
 - CREATE TABLE Member
 - (uid INTEGER NOT NULL,
 - CHECK(uid IN (SELECT uid FROM User)), ...);
 - Is it a referential integrity constraint?
 - Not quite; not checked when the User is modified

Review of Data Manipulation Language

- **SELECT A₁, A₂, ..., A_n**
FROM R₁, R₂, ..., R_m
WHERE condition;
- Also called as SPJ (selection-projection-join) query
- Can correspond to (but not exactly equivalent to) relational algebra query:

$$\pi_{A_1, A_2, \dots, A_n}(\sigma_{\text{condition}}(R_1 \times R_2 \times \dots \times R_m))$$

Why SFW statements

- Out of many possible ways of structuring SQL statements, *why did the designers choose SELECT-FROM-WHERE?*
 - A large number of queries can be written using only selection, projection, and cross-product (or join)
 - Any query that uses only these operators can be written in a canonical form:
$$\pi_L \left(\sigma_p (R_1 \times \dots \times R_m) \right)$$
 - SELECT-FROM-WHERE captures this canonical form

Example Database

User

<i>uid</i>	<i>name</i>	<i>age</i>	<i>pop</i>
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

Group

<i>gid</i>	<i>name</i>
abc	Book Club
gov	Student Government
dps	Dead Putting Society
...	...

Member

<i>uid</i>	<i>gid</i>
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

Example: reading a table

- **SELECT * FROM User;**
 - Single-table query, so no cross product here
 - **WHERE** clause is optional
 - * is a short hand for “all columns”

Example: selection and projection

- Name of users under 18
 - `SELECT name
FROM User
WHERE age < 18`
- When was Lisa born?
 - `SELECT 2025-age
FROM User
WHERE name = 'Lisa';`
 - SELECT list can contain expressions
 - Can also use built-in functions such as SUBSTR, ABS, etc.
 - String literals (case sensitive) are enclosed in single quotes

Example: join

- ID's and names of groups with a user whose name contains “Simpson”

User

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3

Group

gid	name
abc	Book Club
gov	Student Government
dps	Dead Putting Society
...	...

Member

uid	gid
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

Example: join

- ID's and names of groups with a user whose name contains "Simpson"
 - `SELECT Group.gid, Group.name
FROM User, Member, Group
WHERE User.uid = Member.uid
AND Member.gid = Group.gid
AND User.name LIKE '%Simpson%';`
 - **LIKE** matches a string against a pattern
 - % matches any sequence of zero or more characters
 - Okay to omit table_name in table_name.column_name if column_name is unique

User

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

Member

uid	gid
142	dps
123	gov
857	abc
857	gov
456	abc
456	gov
...	...

Group

gid	name
abc	Book Club
gov	Student Government
dps	Dead Putting Society
...	...

SQL Intermediate

Today's Agenda

- Data Definition Language
- **Data Manipulation Language**
 - Basic Queries (SELECT-FROM-WHERE)
 - **ORDER BY**

Order By

- SELECT ... FROM ... WHERE ... ORDER BY output_column [ASC|DESC], ...;
(In the future: SELECT [DISTINCT] ... FROM ... WHERE ... GROUP BY ...
HAVING ... ORDER BY output_column [ASC|DESC], ...;)
- ASC = ascending, DESC = descending
- **Semantics:** After the SELECT list has been computed and optional
duplicate elimination has been carried out, ***sort the output according to
ORDER BY specification***

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3
...

Order By example

- List all users, sort them by popularity (descending) and name (ascending)
 - SELECT uid, name, age, pop FROM User
ORDER BY pop DESC, name;
 - ASC is the default option
 - Strictly speaking, only output columns can appear in the ORDER BY clause (although some DBMS support more)
 - Can use sequence numbers instead of names to refer to output columns: **ORDER BY 4 DESC, 2;**

Today's Agenda

- Data Definition Language
- **Data Manipulation Language**
 - Basic Queries (SELECT-FROM-WHERE)
 - ORDER BY
 - **Set Operations**

Forcing set semantics

- IDs of all pairs of users that belong to one group

- `SELECT m1.uid AS uid1, m2.uid AS uid2
FROM Member AS m1, Member AS m2
WHERE m1.gid = m2.gid AND m1.uid > m2.uid;`

(finds pairs of members who belong to the same group.)

- Say Lisa and Ralph are in both the book club and the student government*

*Forcing set semantics means ensuring that a query returns **unique results** by explicitly removing duplicates, typically using operations like `DISTINCT`, rather than allowing duplicate entries in the result set*

- `SELECT DISTINCT m1.uid AS uid1, m2.uid AS uid2 ...`

- With `DISTINCT`, all duplicate (uid1, uid2) pairs are removed from the output

User				Member	
uid	name	age	pop	uid	gid
142	Bart	10	0.9	142	dps
123	Milhouse	10	0.2	123	gov
857	Lisa	8	0.7	857	abc
456	Ralph	8	0.3	857	gov
...	456	abc

Group	
gid	name
abc	Book Club
gov	Student Government
dps	Dead Putting Society
...	...

Set Semantics of SFW

- **SELECT [DISTINCT] E₁, E₂, ..., E_n**
FROM R₁, R₂, ..., R_m
WHERE condition;

SQL set

- **UNION**, **EXCEPT**, **INTERSECT**
 - Set semantics
 - Duplicates in input tables, if any, are first eliminated
 - Duplicates in result are also eliminated (for UNION)
 - Exactly like set **U**, **-**, and **n** in relational algebra

Examples of set operations

- Poke (uid1, uid2, timestamp)
 - (SELECT uid1 FROM Poke)
EXCEPT
(SELECT uid2 FROM Poke);

Examples of set operations

- Poke (uid1, uid2, timestamp)
 - (SELECT uid1 FROM Poke)
EXCEPT
(SELECT uid2 FROM Poke);
 - **Users who poked others but never got poked by others**
- The EXCEPT operator returns the **set difference** between the two subqueries.
- It gives you all the uid1 values (users who poked others) that are **not** present in the uid2 values (users whom someone poked).

Today's Agenda

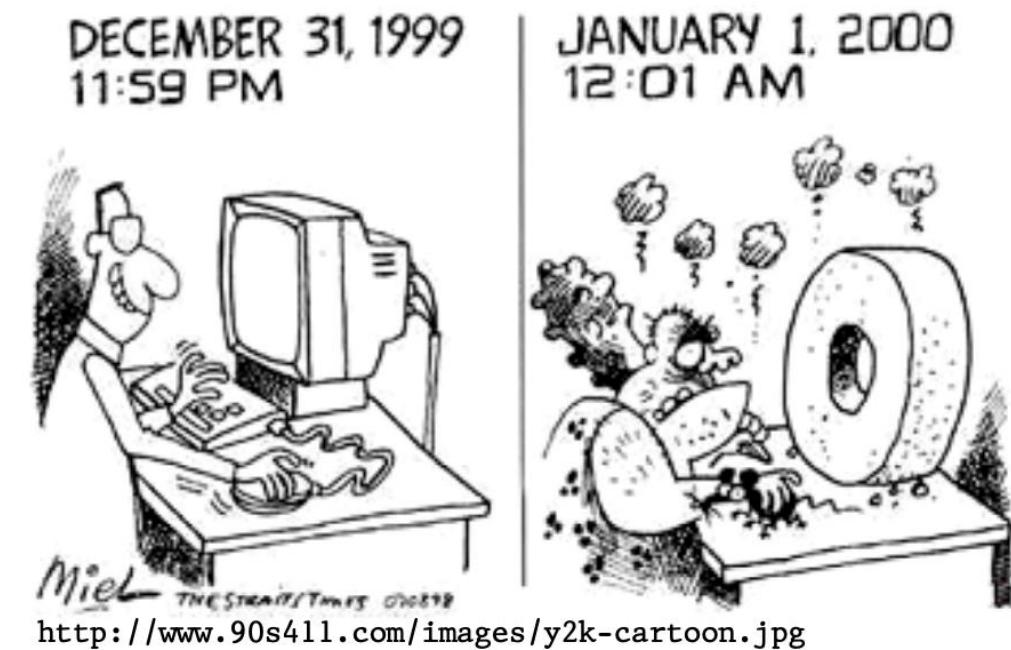
- Data Definition Language
- **Data Manipulation Language**
 - Basic Queries (SELECT-FROM-WHERE)
 - ORDER BY
 - Set Operations
 - **Null Values**

Missing Information

- Example: User (uid, name, age, pop)
- Value **unknown**
 - We do not know Nelson's age
- Value **not applicable**
 - Suppose pop is based on interactions with others on our social networking site
 - Nelson is new to our site; what is his pop?

Solution 1

- Dedicate a value from each domain (type)
 - pop cannot be -1, so use -1 as a special value to indicate a missing or invalid pop
- Leads to incorrect answers if not careful
 - SELECT AVG(pop) FROM User;
- Complicates applications
 - SELECT AVG(pop) FROM User WHERE pop <> -1;
- Perhaps the value is not as special as you think!
 - Ever heard of the Y2K bug? "00" was used as a missing or invalid year value



<http://www.90s411.com/images/y2k-cartoon.jpg>

Solution 2 (using a flag)

- A **valid-bit** for every column
- User (uid, name,
 name_is_valid ,
 age, age_is_valid ,
 pop, pop_is_valid)
- Complicates schema and queries
- SELECT AVG(pop) FROM User WHERE pop_is_valid ;

Solution 3

- **Decompose the table; missing row = missing value**
 - UserName (uid, name)
 - UserAge (uid, age)
 - UserPop (uid, pop)
 - UserID (uid)
 - Still complicates schema and queries
 - How to get all information about users in a table?
 - Natural join doesn't work!

SQL's solution

- A special value **NULL**
 - For every domain
 - Special rules for dealing with NULL's
- Example: User (uid, name, age, pop)
 - <789, "Nelson", NULL, NULL>

Computing with NULLs

- When we operate on a NULL and another value (including another NULL) using +, −, etc., the result is NULL
- Aggregate functions ignore NULL, except COUNT(*) (since it counts rows)

IS NULL/IS NOT NULL

- Example: Who has NULL pop values?
 - `SELECT * FROM User WHERE pop = NULL;`
 - Does not work; never returns anything
- SQL introduced special, built-in predicates **IS NULL** and **IS NOT NULL**
 - `SELECT * FROM User WHERE pop IS NULL;`

Outerjoin motivation

First let's see: what is the natural join result?

<i>Employee</i>			<i>Dept</i>	
Name	Empld	DeptName	DeptName	Manager
Harry	3415	Finance	Sales	Harriet
Sally	2241	Sales	Production	Charles
George	3401	Finance		
Harriet	2202	Sales		
Tim	1123	Executive		

Outerjoin motivation

First let's see: what is the natural join result?

Employee			Dept					
Name	Empld	DeptName	DeptName	Manager	Name	Empld	DeptName	Manager
Harry	3415	Finance	Sales	Harriet				
Sally	2241	Sales	Production	Charles				
George	3401	Finance						
Harriet	2202	Sales						
Tim	1123	Executive						

⊗ =

Outerjoin motivation

First let's see: what is the natural join result?

Employee			Dept					
Name	Empld	DeptName	DeptName	Manager	Name	Empld	DeptName	Manager
Harry	3415	Finance	Sales	Harriet	Sally	2241	Sales	Harriet
Sally	2241	Sales	Production	Charles	Harriet	2202	Sales	Harriet
George	3401	Finance						
Harriet	2202	Sales						
Tim	1123	Executive						

Outerjoin motivation

First let's see: what is the natural join result?

Employee			Dept					
Name	Empld	DeptName	DeptName	Manager	Name	Empld	DeptName	Manager
Harry	3415	Finance	Sales	Harriet	Sally	2241	Sales	Harriet
Sally	2241	Sales	Production	Charles	Harriet	2202	Sales	Harriet
George	3401	Finance						
Harriet	2202	Sales						
Tim	1123	Executive						

Bad, other employee information get lost, because their departments do not have a record in the Dept table

Optional: Left Outer Natural Join

- In a left outer join, Employee rows without a matching Department row appear in the result, but not vice versa.

<i>Employee</i>		
Name	Empld	DeptName
Harry	3415	Finance
Sally	2241	Sales
George	3401	Finance
Harriet	2202	Sales
Tim	1123	Executive

<i>Dept</i>	
DeptName	Manager
Sales	Harriet
Production	Charles

<i>Employee ⚫ Dept</i>			
Name	Empld	DeptName	Manager
Harry	3415	Finance	ω
Sally	2241	Sales	Harriet
George	3401	Finance	ω
Harriet	2202	Sales	Harriet
Tim	1123	Executive	ω

Optional: Left Outer Natural Join

- \bowtie

$$(R \bowtie S) \cup ((R - \pi_{r_1, r_2, \dots, r_n}(R \bowtie S)) \times \{(\omega, \dots, \omega)\})$$

- In a left outer join, Employee rows without a matching Department row appear in the result, but not vice versa.

<i>Employee</i>		
Name	Empld	DeptName
Harry	3415	Finance
Sally	2241	Sales
George	3401	Finance
Harriet	2202	Sales
Tim	1123	Executive

<i>Dept</i>	
DeptName	Manager
Sales	Harriet
Production	Charles

<i>Employee \bowtie Dept</i>			
Name	Empld	DeptName	Manager
Harry	3415	Finance	ω
Sally	2241	Sales	Harriet
George	3401	Finance	ω
Harriet	2202	Sales	Harriet
Tim	1123	Executive	ω

SQL Continued..

- Data Definition Language
- **Data Manipulation Language**
 - Basic Queries (SELECT-FROM-WHERE)
 - ORDER BY
 - Set Operations
 - Null Values
 - **Aggregation**

Aggregates

- Standard SQL aggregate functions: **COUNT, SUM, AVG, MIN, MAX**
- Example: number of users under 18, and their average popularity
 - `SELECT COUNT(*)
FROM User
WHERE age < 18;`
 - **COUNT(*)** counts the number of rows

Aggregates with Distinct

- Example: How many users are in some group?
- `SELECT COUNT(DISTINCT uid)`
 `FROM Member;`
is equivalent to:
• `SELECT COUNT(*)`
 `FROM (SELECT DISTINCT uid FROM Member);`

Grouping

- `SELECT ... FROM ... WHERE ...
GROUP BY list_of_columns`
- Example: compute average popularity for each age group
`SELECT age, AVG(pop)
FROM User
GROUP BY age`

Example of Grouping By

- SELECT age, AVG(pop) FROM User GROUP BY age;

uid	name	age	pop
142	Bart	10	0.9
857	Lisa	8	0.7
123	Milhouse	10	0.2
456	Ralph	8	0.3

Compute GROUP BY: group rows according to the values of GROUP BY columns

The diagram illustrates the process of performing a GROUP BY operation. It starts with a table of user data (uid, name, age, pop). An arrow points to a second table where rows are grouped by age (uid, name, pop for age 10 are highlighted in yellow; uid, name, pop for age 8 are highlighted in green). A final arrow points to a third table showing the results of the SELECT query, which includes the average population for each age group.

uid	name	age	pop
142	Bart	10	0.9
123	Milhouse	10	0.2
857	Lisa	8	0.7
456	Ralph	8	0.3

Compute SELECT
for each group

age	avg_pop
10	0.55
8	0.50

Example of Aggregates (with no Group By)

- An aggregate query with no GROUP BY clause = all rows go into one group

```
SELECT AVG(pop) AS avg_pop  
FROM User
```

uid	name	age	pop
142	Bart	10	0.9
857	Lisa	8	0.7
123	Milhouse	10	0.2
456	Ralph	8	0.3

Group all rows
into one group

Aggregate over
the whole group

uid	name	age	pop	avg_pop
142	Bart	10	0.9	
857	Lisa	8	0.7	
123	Milhouse	10	0.2	
456	Ralph	8	0.3	0.525

Having

- Used to filter groups based on the group properties (e.g., aggregate values, GROUP BY column values)
- SELECT ... FROM ... WHERE ... GROUP BY ...
HAVING condition;

Having examples

- List the average popularity for each age group with more than a hundred users
 - ```
SELECT age, AVG(pop)
 FROM User
 GROUP BY age
 HAVING COUNT(*) > 100;
```

    - Can be written using WHERE and table subqueries
  - Find average popularity for each age group over 10
    - ```
SELECT age, AVG(pop)
  FROM User
 GROUP BY age
 HAVING age > 10;
```

 - Can be written using WHERE without table subqueries